

Short-Term Conflict Resolution for Unmanned Aircraft Traffic Management

Hao Yi Ong and Mykel J. Kochenderfer

Stanford University

September 22, 2015

Outline

Introduction

Mathematical formulation

Approach

Numerical experiments

UTM- α : A distributed framework

Conclusion and future work

Unmanned aircraft traffic management


- ▶ automating conflict avoidance is critical to integrating small unmanned aerial systems (UAS) to civil airspace
- ▶ automation to augment controllers



Conflict avoidance

given potential conflicts between drones, find the best set of advisories

- ▶ focus on horizontal or co-altitude conflict resolution
 - NASA's proposed airspace for UAS is under 150 m (500 ft)
 - flight altitudes need to avoid disruption and buildings
- ▶ conflict defined as loss of minimum separation between aircraft



Complications

- ▶ multiagent problem
 - system must coordinate between many aircraft
 - large search space for solution
- ▶ uncertainty in system
 - imperfect sensor measurements of current state
 - variable pilot response, vehicle performance, etc. affect future path
- ▶ appropriate trade-off between safety and efficiency not obvious

Goals

robust and efficient method


- ▶ tractable approach to complex stochastic problem
 - account for uncertainty in large problem with reasonable compute resources
- ▶ balances between airspace safety and efficiency
 - ensure safety and provide timely conflict alerts to aircraft
- ▶ arbitrary-scale optimization
 - real-time conflict resolution for large airspace

Previous approaches

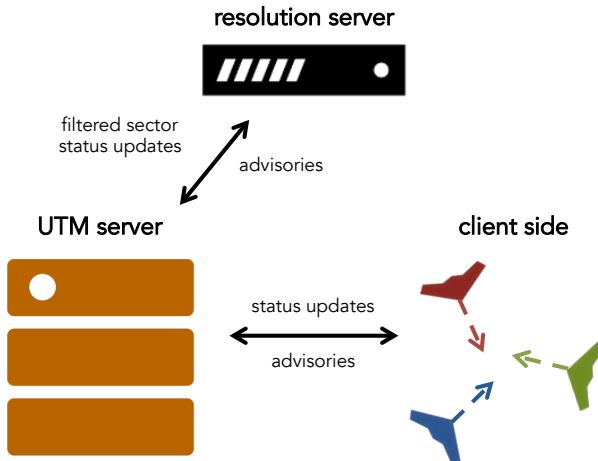
- ▶ mathematical programming (MIP, SCP) [SVFH05, ASD12]
 - can work well for simple vehicle networks
- ▶ distributed convex optimization [OG15]
 - hard to incorporate stochastic objectives/constraints
- ▶ Markov decision process formulation (ACAS X) [KC11, CK12]
 - assumes “white noise” accelerations for intruder aircraft
- ▶ and many more... [KY00]

Overview

idea: decomposition + coordination



Conflict Resolution as a UTM service



Conflict Resolution as a Standalone service

- ▶ derivative UTM client service
 - subscribes to UTM server to track aircraft and deconflict routes
 - pub-sub system that operators subscribe to to receive advisories
- ▶ current implementation uses standalone model
 - still unclear about end architecture for resolution service
 - clean approach to get started with prototype system
 - implemented with scalability and modularity in mind

Outline

Introduction

Mathematical formulation

Approach

Numerical experiments


UTM- α : A distributed framework

Conclusion and future work


Markov decision process (MDP)

defined by the tuple $(\mathcal{S}, \mathcal{A}, T, R)$

- ▶ \mathcal{S} and \mathcal{A} are the sets of all possible states and actions, respectively
- ▶ $T(s, a, s')$ gives the probability of transitioning into state s' by taking action a at the current state s
- ▶ $R(s, a)$ gives the reward for taking action a at the current state s



Value iteration



- ▶ want to find the optimal policy $\pi^*(s)$
 - gives action that maximizes the utility $Q^*(s, a)$ from any given state

$$\pi^*(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a)$$

- ▶ value iteration updates value function guess \hat{Q} until convergence
 - expensive one-off compute but cheap policy extraction (\hat{Q} lookup)

$$\hat{Q}(s, a) := R(s, a) + \sum_{s' \in \mathcal{S}} T(s, a, s') \max_{a' \in \mathcal{A}} \hat{Q}(s', a')$$


Multiagent MDP

extension of MDP to cooperative multiagent setting

- ▶ similar to case where centralized planner has access to system state
- ▶ also defined by the tuple $(\mathcal{S}, \mathcal{A}, T, R)$, except
 - \mathcal{S} and \mathcal{A} are all possible joint states and actions
 - T and R operate on elements of \mathcal{S} and \mathcal{A}

Short-term conflict resolution

- ▶ horizontal conflict resolution between n aircraft
 - conflict defined as loss of minimum separation distance, 500 m
 - aircraft at risk if it could experience a conflict within two minutes
- ▶ alert aircraft that need corrective maneuvers
 - but not to the extent that alerts become a nuisance



Resolution advisories

- ▶ at each time step $T = 5$ s, system issues a joint advisory
 - joint advisory ϕ chosen out of a finite set of corrective bank angles for each aircraft
- ▶ action set $\mathcal{A} = \{-20^\circ, -10^\circ, 0^\circ, 10^\circ, 20^\circ, \text{COC}\}^n$
 - positive and negative angles correspond to left and right banks
 - COC is a clear-of-conflict status advisory—nothing needs to be done
- ▶ higher resolution comes at the cost of computational complexity

Dynamics


► i th aircraft state $s_i = (x_i, y_i, \psi_i, v_i, p_i)$

- latitude and longitude x_i and y_i
- heading ψ_i
- groundspeed v_i
- responsiveness indicator p_i

► aircraft follow Dubin's kinematic model


- simplicity avoids risk of overfitting complicated models

$$\dot{x}_i = v_i \cos \psi_i \quad \dot{y}_i = v_i \sin \psi_i \quad \dot{\psi}_i = \frac{g \tan \phi_i}{v_i}$$



Pilot model

- ▶ advisory response determined stochastically by Bernoulli process
- ▶ 5 s average response delay to first corrective advisory as recommended by ICAO [ICA07]
 - when responding, the pilot executes the advisory for 5 s time step





Reward function

balance competing objectives

► maximize safety

► minimize disruption



Outline

Introduction

Mathematical formulation

Approach


Numerical experiments

UTM- α : A distributed framework


Conclusion and future work

Curse of dimensionality


- ▶ in an MMDP, \mathcal{S} and \mathcal{A} are all possible joint states and actions
- ▶ **problem:** \mathcal{S} and \mathcal{A} scale exponentially with number of agents



Decompose and coordinate



MMDP decomposition




Pairwise conflict resolution

decompose full MMDP into $O(n^2)$ pairwise encounters

- ▶ action set, pilot response model, and reward function unchanged
- ▶ use relative positions and bearing to reduce state space
 - arbitrarily set one aircraft as ownship, and the other as intruder
 - speeds remain in global reference frame

$$s = (x^{\text{rel}}, y^{\text{rel}}, \psi^{\text{rel}}, v^{\text{ownship}}, v^{\text{intruder}}, p^{\text{ownship}}, p^{\text{intruder}})$$



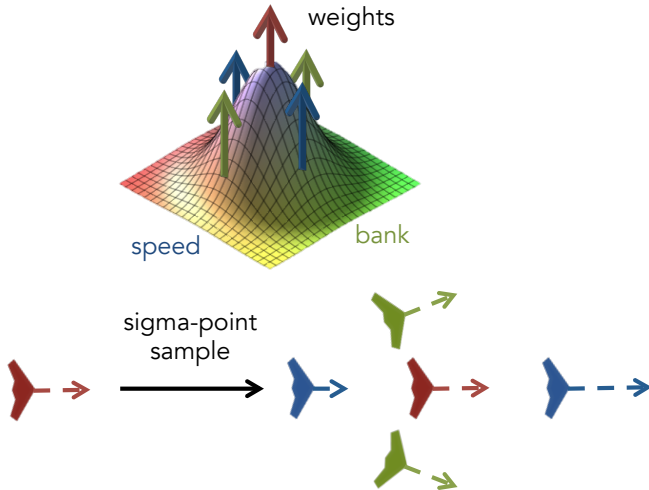
Discretization

- ▶ in general, no analytical solution to problem with continuous state space
- ▶ approximate value function \hat{Q} over discretized state space
 - discretize state space via multilinear interpolation
 - iteratively update \hat{Q} until convergence

$$\hat{Q}(s, a) := R(s, a) + \sum_{s' \in \mathcal{S}} T(s, a, s') \max_{a' \in \mathcal{A}} \hat{Q}(s', a')$$

Approximating environment noise

sigma-point sampling



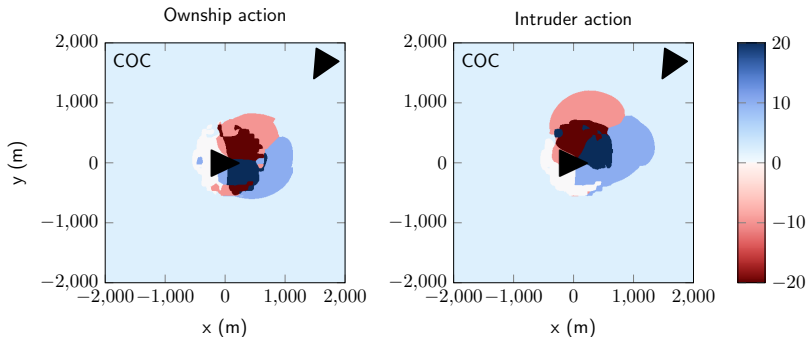
Implementation

- ▶ discretize continuous state space into 9.6 million states
 - expensive but one-off computation
- ▶ \hat{Q} converges in 7 hours (40 iterations over $\mathcal{S} \times \mathcal{A}$)
 - solver: parallel implementation in Julia
 - machine: 20 2.3 GHz Intel Xeon cores with 125 GB RAM

Variable	Minimum	Maximum	Number of values
$x^{\text{rel}}, y^{\text{rel}}$	−3000 m	3000 m	51
ψ^{rel}	0°	360°	37
$v^{\text{ownship}}, v^{\text{intruder}}$	10 m s ^{−1}	20 m s ^{−1}	5
$p^{\text{ownship}}, p^{\text{intruder}}$	–	–	2

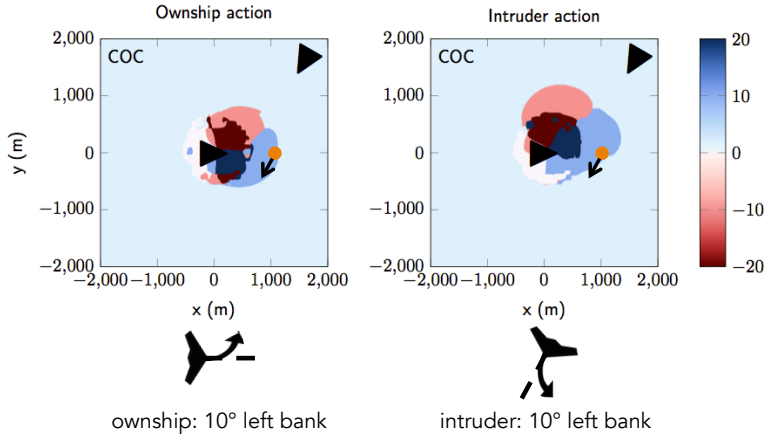
Pairwise policy visualization

passing intrusion (240° bearing)



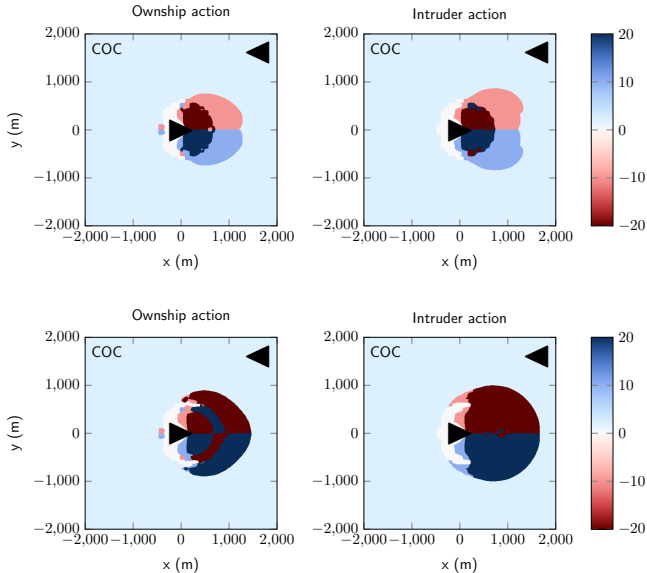
Example advisory execution

passing intrusion at (1000 m, 0 m)



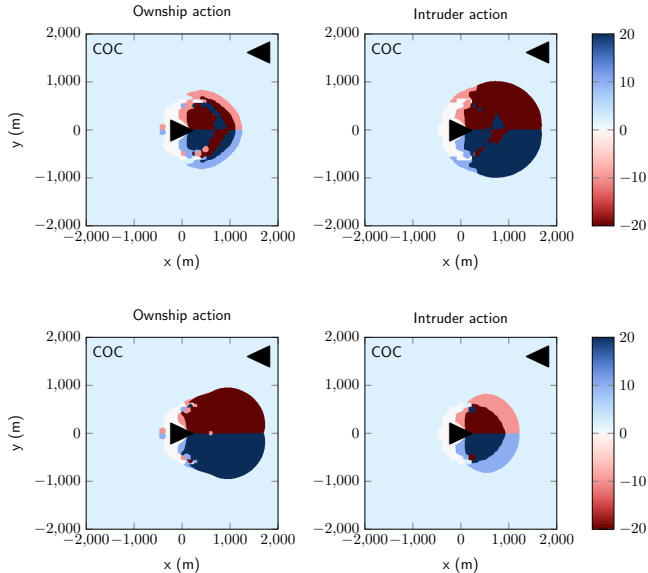
Accounting for pilot response

pilots responsive (top) and non-responsive (bottom)




Accounting for pilot response

ownship responsive (top) and intruder responsive (bottom)



Multithreat coordination



Utility fusion

idea: combine pairwise utilities Q_{ij} to form proxy for global utility Q

- ▶ Q_{ij} operates on the pairwise MDP state-action pair (s_{ij}, a_{ij}) for aircraft i and j
- ▶ fusion strategies $\pi(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$

- max-min: “worst-case”

$$Q^{\min}(s, a) = \min_{i < j} \{U_{ij}(s_{ij}, a_{ij})\}$$


- max-sum: “average”

$$Q^{\text{sum}}(s, a) = \sum_{i < j} U_{ij}(s_{ij}, a_{ij})$$

problem: action space \mathcal{A} scales exponentially with number of aircraft


Search heuristic: Alternating maximization

- ▶ maximize global utility by varying one action and fixing the rest
- ▶ e.g., fix **drone 1** and **2's** actions (**R** and **L**) and vary **drone 3's** action




Distributed variant I: Parallel search

greedy local maximization




repeat until  convergence

broadcast action




Distributed variant II: Consensus search


parallel alt. maximization




greedy local maximization



broadcast joint action




incorporate msgs



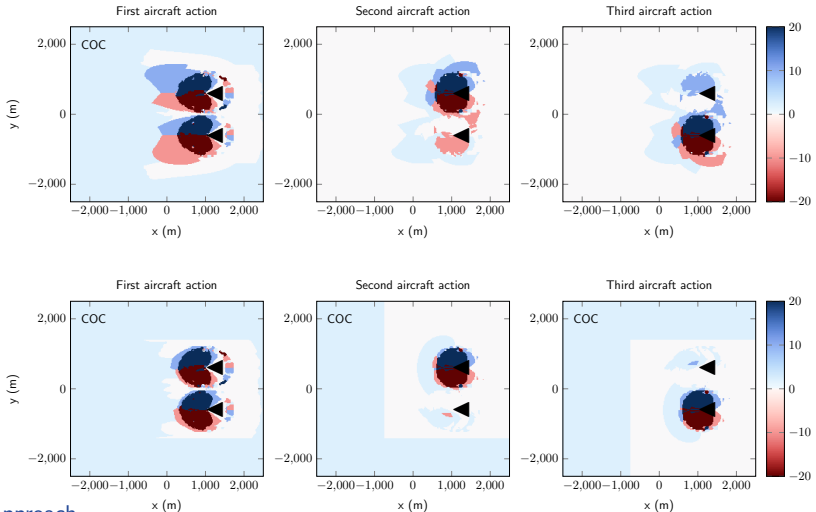
Consensus search message

- ▶ parallel search could incur large communication cost
- ▶ messages reduce communication between compute nodes, but risk yielding unsynchronized joint action across aircraft
- ▶ incorporate messages by “averaging” what intruders think ownship should do to approximate consensus



Joint policy visualization

triple aircraft encounter max-min (top) and max-sum policies (bottom)



Outline

Introduction

Mathematical formulation

Approach


Numerical experiments

UTM- α : A distributed framework

Conclusion and future work

Encounter model

- ▶ generate uniformly random starting aircraft states in annulus
 - speeds uniformly random
 - headings initialized to point to annulus center
 - resample if loss of separation



Aircraft model

- ▶ ODE solver for dynamics
 - continuous Dubin's kinematic model
 - Gaussian noise in acceleration and banking
- ▶ corrective bank angles mapped to PID control policies

Baseline methods


- ▶ closest threat
 - multithreat scenario decomposed into pairwise encounters
 - aircraft execute solution to encounter with closest intruder
- ▶ uncoordinated
 - aircraft assumes all other aircraft are white-noise intruders
 - executes greedy local maximization to find own action

Experiment


- ▶ >1 million simulations from 2 to 10 aircraft
 - code: Julia implementation
 - machine: 3.4 GHz Intel i7 processor with 32 GB RAM
- ▶ ~10 ms decision times (vs. 5 s decision period)
 - serial solve times of <10 ms for up to 10 aircraft sufficient for server-based resolution system
 - unoptimized code that doesn't account for communication costs for distributed variants

Safety performance

► max-min utility fusion



► max-sum utility fusion




- distributed variants perform as well as serial version
- coordinated methods >10 times better than closest threat heuristic
- coordinated methods $\sim 10\%$ better than uncoordinated method

Safety performance takeaways

- ▶ no near mid-air collisions (NMACs) in simulations
 - NMAC threshold distance defined at 30 m (100 ft)
 - due to penalty on smaller separation distances even in conflict
- ▶ simulations validate all algorithmic variants
 - distributed variants perform as well as serial version
 - coordinated methods significantly better than baselines

Trade-off: Safety vs. alert rate

- ▶ vary relative penalty between loss of separation and disruption
- ▶ best performance at “knee” of Pareto frontier



Outline

Introduction

Mathematical formulation

Approach

Numerical experiments

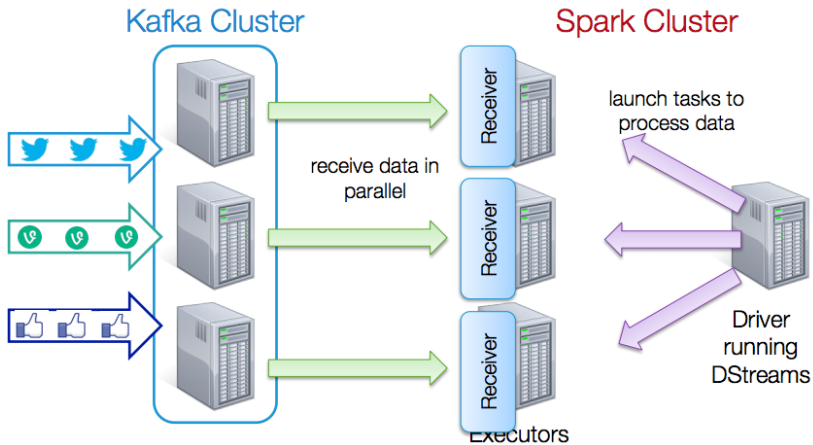
UTM- α : A distributed framework

Conclusion and future work

Practical conflict avoidance

- ▶ recall: pub-sub system that UTM clients subscribe to for advisories
 - standalone system that subscribes to UTM server for flight tracking
- ▶ **goals:** robustness, scalability, and modularity
 - cluster computing framework for large-scale data processing
 - streaming analytics for real-time conflict resolution
 - built-in system fault-tolerance

System design



System design


► why Kafka?

- designed as unified, high-throughput, low-latency platform for handling real-time data feeds
- open source with active industry use (originally developed by LinkedIn)

► why Spark?


- easy, high-level API
- ease of operations with minimal fuss over fault-tolerance
- enormous, active community and industry deployments (CISCO, Netflix, Intel, . . .)

Driver-worker model



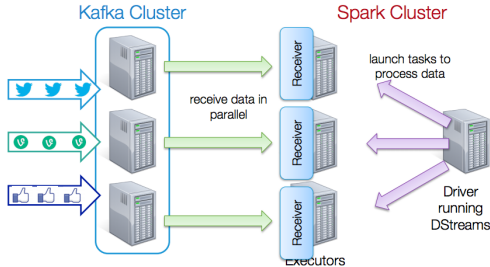
- ▶ driver-worker model for conflict resolution
 - driver loads up policy object that contains lookup table and algorithm for conflict resolution
 - workers receives policy object via broadcast and subsequently delegated tasks by driver
- ▶ formatted conflict data stored as an “infinite” stream of resilient distributed datasets (RDDs)
 - distributed, partitioned collection of conflict objects (JSON)
 - automatically rebuilt on failure

Pub-sub model



- ▶ ingestor publishes conflict objects to conflict Kafka topic
 - subscribes to and ingests UTM server stream and crudely identifies potential conflicts by proximity
 - formats conflict data as JSON strings
- ▶ advisor publishes advisories to advisory Kafka topic
 - producers are worker nodes in cluster generating advisories
 - consumers are UTM operator clients flying drones

Working model



- ▶ Spark driver launches workers/executors and broadcasts policy object
- ▶ workers receive conflict data in parallel from conflict Kafka topic as RDD stream
- ▶ workers publish advisories to advisory Kafka topic for UTM operator clients

Outline

Introduction

Mathematical formulation

Approach

Numerical experiments

UTM- α : A distributed framework

Conclusion and future work

Summary

- ▶ three variants of coordination-based conflict resolution algorithm
 - 1 centralized/serial and 2 distributed variants
- ▶ millisecond solve times for real-time application for multiple aircraft
- ▶ robust distributed system for practical conflict avoidance

Implementation

- ▶ policy generation on Julia scientific programming language

`https://github.com/sisl/ConflictAvoidanceDASC`

- ▶ distributed system with Apache Kafka and Spark Streaming

`https://bitbucket.org/sisl/utm-alpha`

Further research

- ▶ tracking aircraft via belief states
 - use the UTM client server to track reported aircraft states
 - generate local set of belief states for all aircraft to better estimate state uncertainty in real-time
- ▶ best action for drone in case of communication loss
 - figure out the response state of the aircraft via belief state tracking
 - where to go in case of communication loss, what flight profile to take, and what UTM should assume about the drones flight
- ▶ dealing with adversarial flights

Further software development

- ▶ standardize resolution advisory and advisory receipt formats with focus on minimizing message size
- ▶ quantify delay times for sending packets between advisory server and clients in order to model it into the problem

References


- ▶ Federico Augugliaro, Angela P Schoellig, and Raffaello D'Andrea.
Generation of collision-free trajectories for a quadrocopter fleet: A sequential convex programming approach.
In IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS), pages 1917–1922, October 2012.
- ▶ J. P. Chryssanthacopoulos and M. J. Kochenderfer.
Decomposition methods for optimized collision avoidance with multiple threats.
Journal of Guidance, Control, and Dynamics, 35(2):398–405, 2012.
- ▶ International Civil Aviation Organization ICAO.
Surveillance, radar and collision avoidance.
International Standards and Recommended Practices, 4, 2007.

References


- ▶ M.J. Kochenderfer and J.P. Chryssanthacopoulos.
Robust airborne collision avoidance through dynamic programming.
Project Report ATC-371, MIT Lincoln Laboratory, 2011.
- ▶ J. K. Kuchar and L. C. Yang.
A review of conflict detection and resolution modeling methods.
IEEE Trans. Intell. Transp. Syst., 1(4):179–189, 2000.
- ▶ H. Y. Ong and J. C. Gerdes.
Cooperative collision avoidance via proximal message passing.
In *American Control Conference (ACC)*, July 2015.
- ▶ Tom Schouwenaars, Mario Valenti, Eric Feron, and Jonathan How.
Implementation and flight test results of MILP-based UAV guidance.
In *IEEE Aerospace Conference*, pages 1–13. IEEE, 2005.

Impact of search timeout

conflict decreases with iteration limit but with diminishing returns



Restriction messages for consensus search



- ▶ instead of broadcasting joint action, aircraft broadcast restriction messages that indicate what other aircraft shouldn't do
 - e.g., reward straight path when it receives no left and no right messages at a decision period